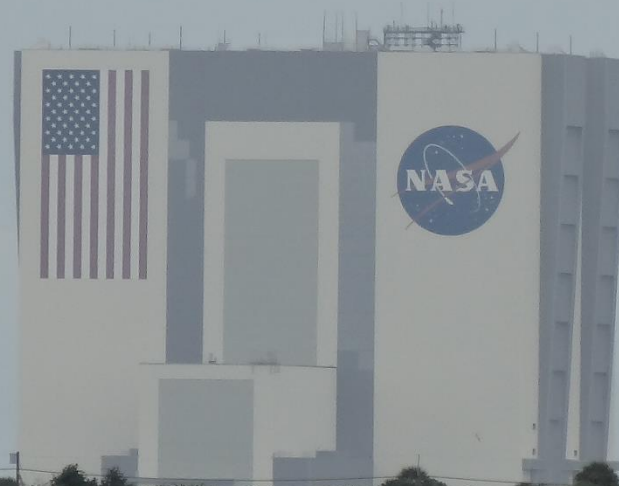


# Deep Learning in Python



## Autoencoders and Pix2pix networks

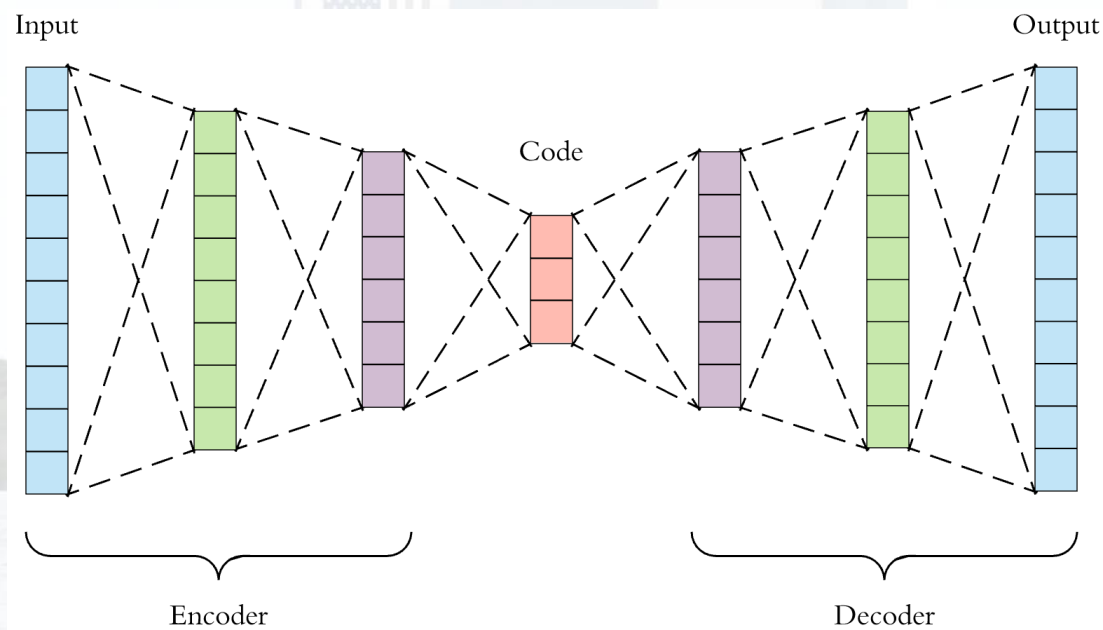
Paweł Kasprowski, PhD, DSc.

# Picture to picture

- The network that generates an image
- Training:
  - input\_image -> network -> output\_image
- Problem:
  - the network should generalize

# Autoencoders

- The network that consists of:
  - Encoder – converts an image into a vector (code)
  - Decoder – converts the code into an image



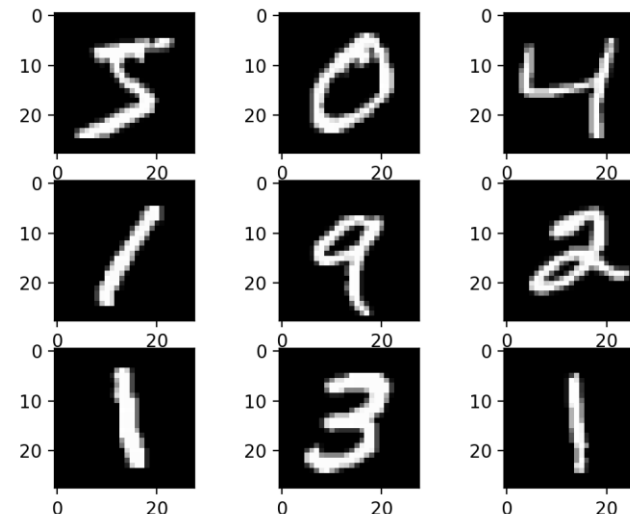
<https://towardsdatascience.com/applied-deep-learning-part-3-autoencoders-1c083af4d798>

# Encoder

- The sample is recalculated to lower dimension
- For instance:
  - image (200x200x3) is encoded to the vector (100)
- The idea:
  - this compressed (latent) representation preserves **the most important** properties of the original object
  - it will be possible to reconstruct **the same** object from the latent representation

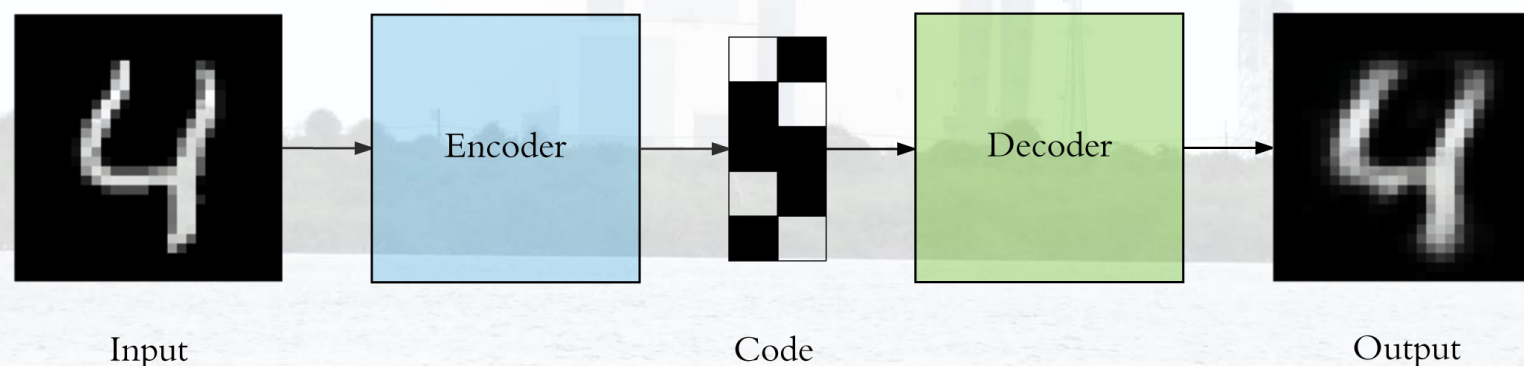
# MNIST dataset

- Handwritten digits
- 10 classes
- 60,000 training examples
- 10,000 test examples
- size: 28x28x1



# The idea

- Encode to the latent vector of size=code\_size
- Decode to the original image
- Training: the same image as input and output!



<https://towardsdatascience.com/applied-deep-learning-part-3-autoencoders-1c083af4d798>



# The simplest autoencoder

## *autoencoders1.ipynb*

- Simple dense network (gets flattened images):

```
code_size = 5
input_img = Input(shape=(28*28,))
code = Dense(code_size, activation='relu')(input_img)
output_img = Dense(28*28, activation='sigmoid')(code)
autoencoder = Model(input_img, output_img)
```
- Training:
  - `autoencoder.compile(optimizer='adam', loss='binary_crossentropy')`
  - `autoencoder.fit(trainSamples, trainSamples, epochs=5)`

# Testing the network

- For different code\_size: 1, 5, 10, 100
- For more sophisticated architecture with two hidden layers:

```
input_img = Input(shape=(input_size,))  
hidden_1 = Dense(128, activation='relu')(input_img)  
code = Dense(code_size, activation='relu')(hidden_1)  
hidden_2 = Dense(128, activation='relu')(code)  
output_img = Dense(input_size, activation='sigmoid')(hidden_2)
```



# Denoising autoencoder

## *autoencoders2.ipynb*

- Creating noisy samples:

```
noise_factor = 0.4
```

```
trainSamples_noisy = trainSamples +  
    noise_factor * np.random.normal(size=trainSamples.shape)
```

```
trainSamples_noisy = np.clip(trainSamples_noisy, 0.0, 1.0)
```

- Train using noisy samples:
  - `autoencoder.fit(trainSamples_noisy, trainSamples, epochs=5)`
- Works better for noisy than for sharp!

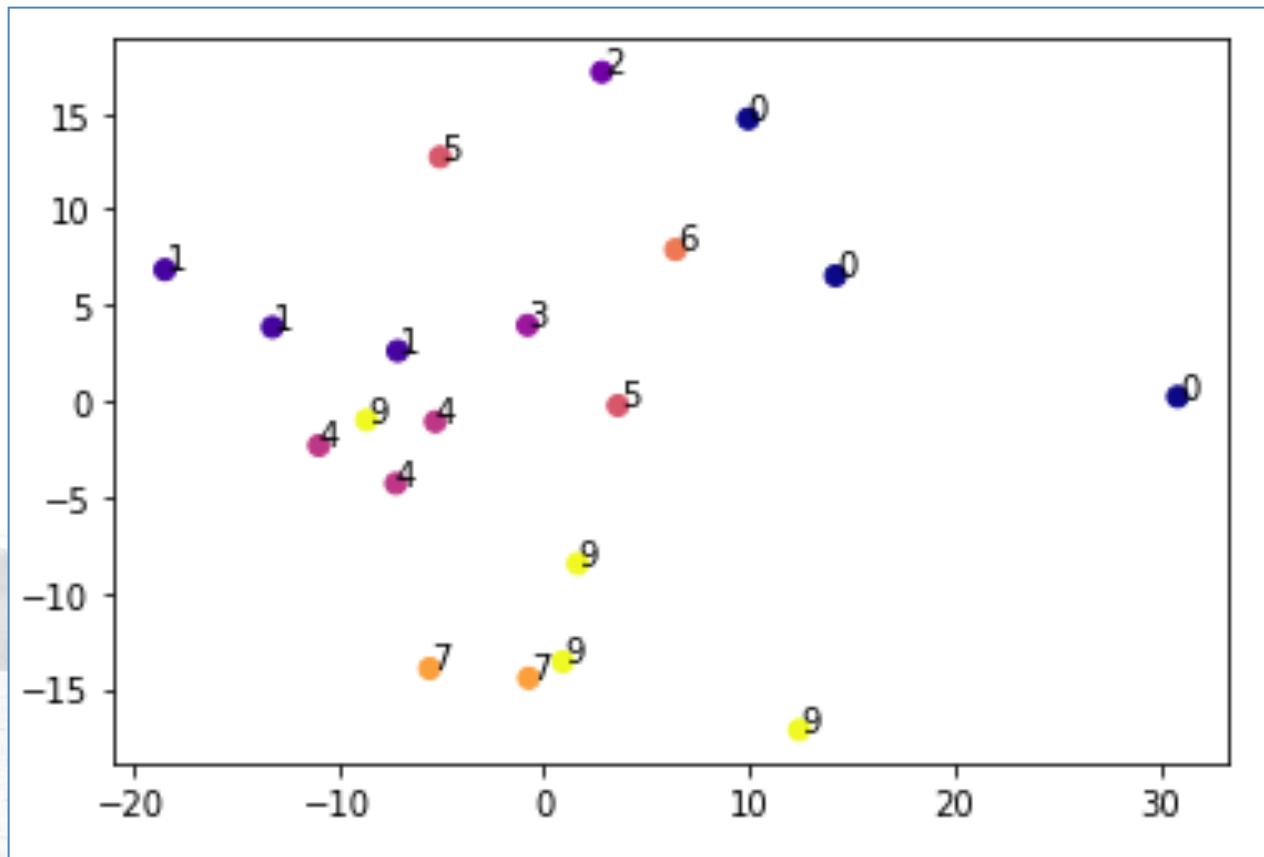
# Analysing codes

- Codes for the same digits should be similar
- Let's map the codes to 2D and analyse on plot
- We will use Principal Component Analysis (PCA) to map it

```
from sklearn.decomposition import PCA
encoder = Model(input_img, code)
vectors = encoder.predict(testSamples[:20]) # get codes
pca = PCA(n_components=2)
vectors2D = pca.fit_transform(vectors) # transform to 2D
plt.scatter(vectors2D[:,0],vectors2D[:,1], c=testLabels[:points])
for i,w in enumerate(vectors):
    plt.annotate(testLabels[i],(vectors2D[i,0],vectors2D[i,1]))
```

# Results

- For code\_size=32 digits in similar places



# Classification of encoded vectors

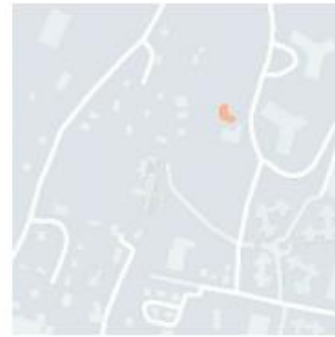
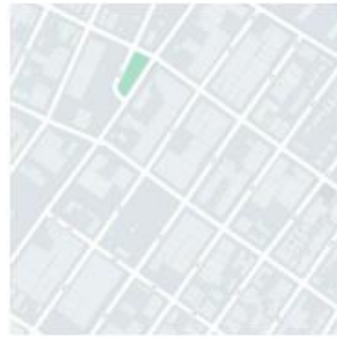
- Prepare vectors:  
    `encoder = Model(input_img, code)`  
    `testVectors = encoder.predict(testSamples)`  
    `trainVectors = encoder.predict(trainSamples)`
- Use kNN to classify vectors  
    `knn_model = KNeighborsClassifier()`  
    `knn_model.fit(trainVectors, trainLabels)`  
    `predLabels = knn_model.predict(testVectors)`
- Results:
  - over 97% for each class!

# Autoencoder applications

- Compression
  - but only for specific data
  - typically jpeg algorithm is better...
- Denoising
  - requires examples!
- Providing latent space vector for future analysis
  - for instance for classification using kNN

# Map example

- Input google maps



- Create autoencoders
- Three architectures:
  - Conv2D
  - Upsampling
  - UNET

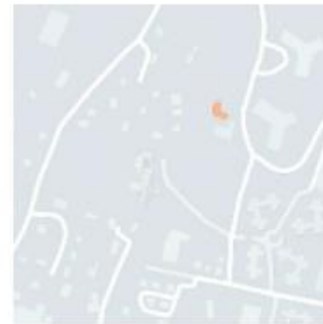
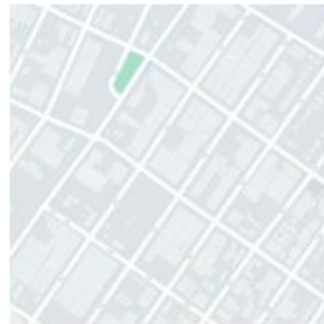
***autoencoder\_models.py***



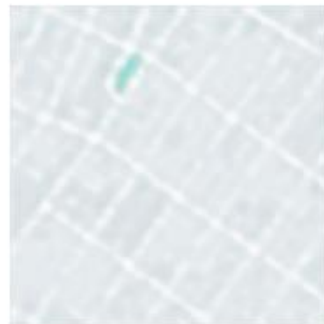
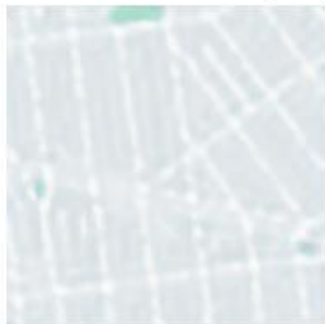
# Upsampling

- Uses UpSampling2D to reconstruct the image
- Result: blurred (***autoencoders\_map\_upsampling.ipynb***)

Original training images



Reconstructed training images



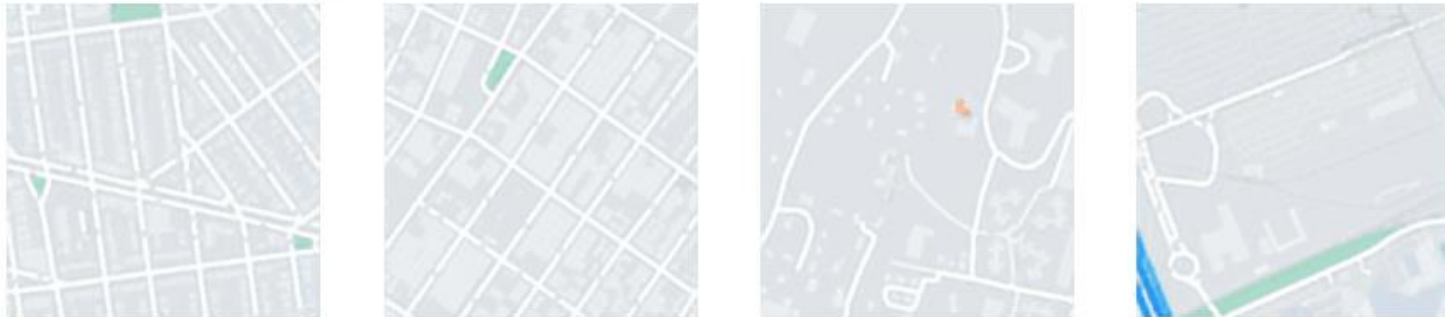
# Conv2DTranspose

- Used Conv2DTranspose layers to decode
- Result: much better and with less iterations  
***autoencoders\_map\_transpose.ipynb***

Original training images

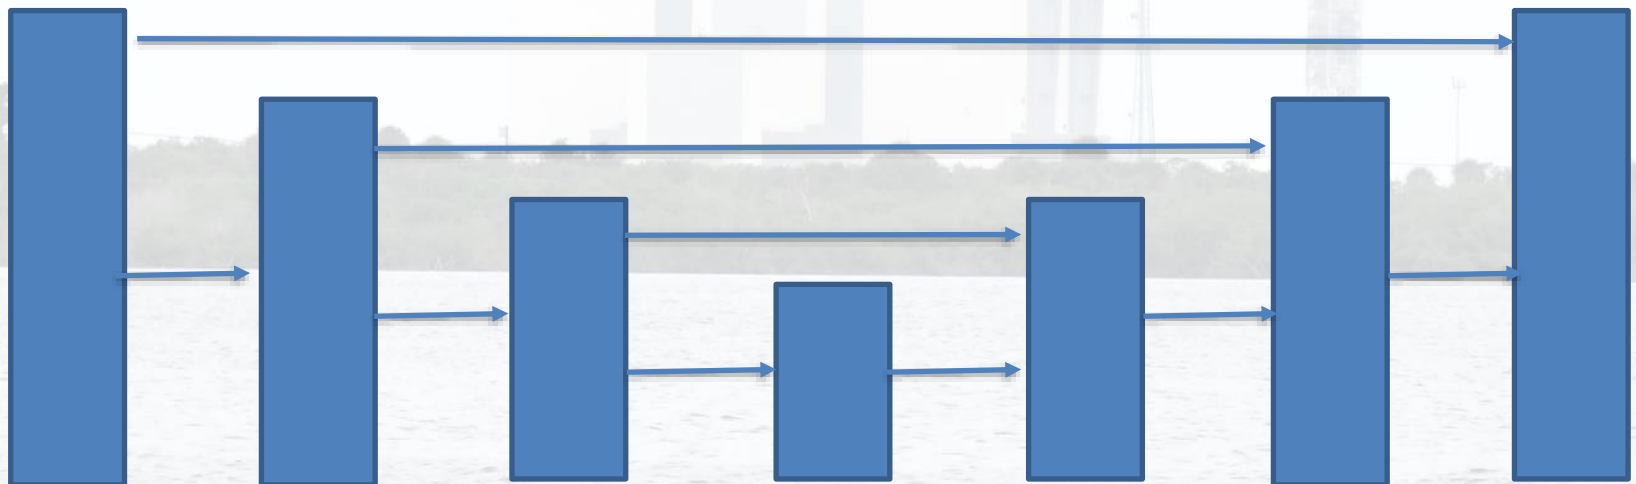


Reconstructed training images

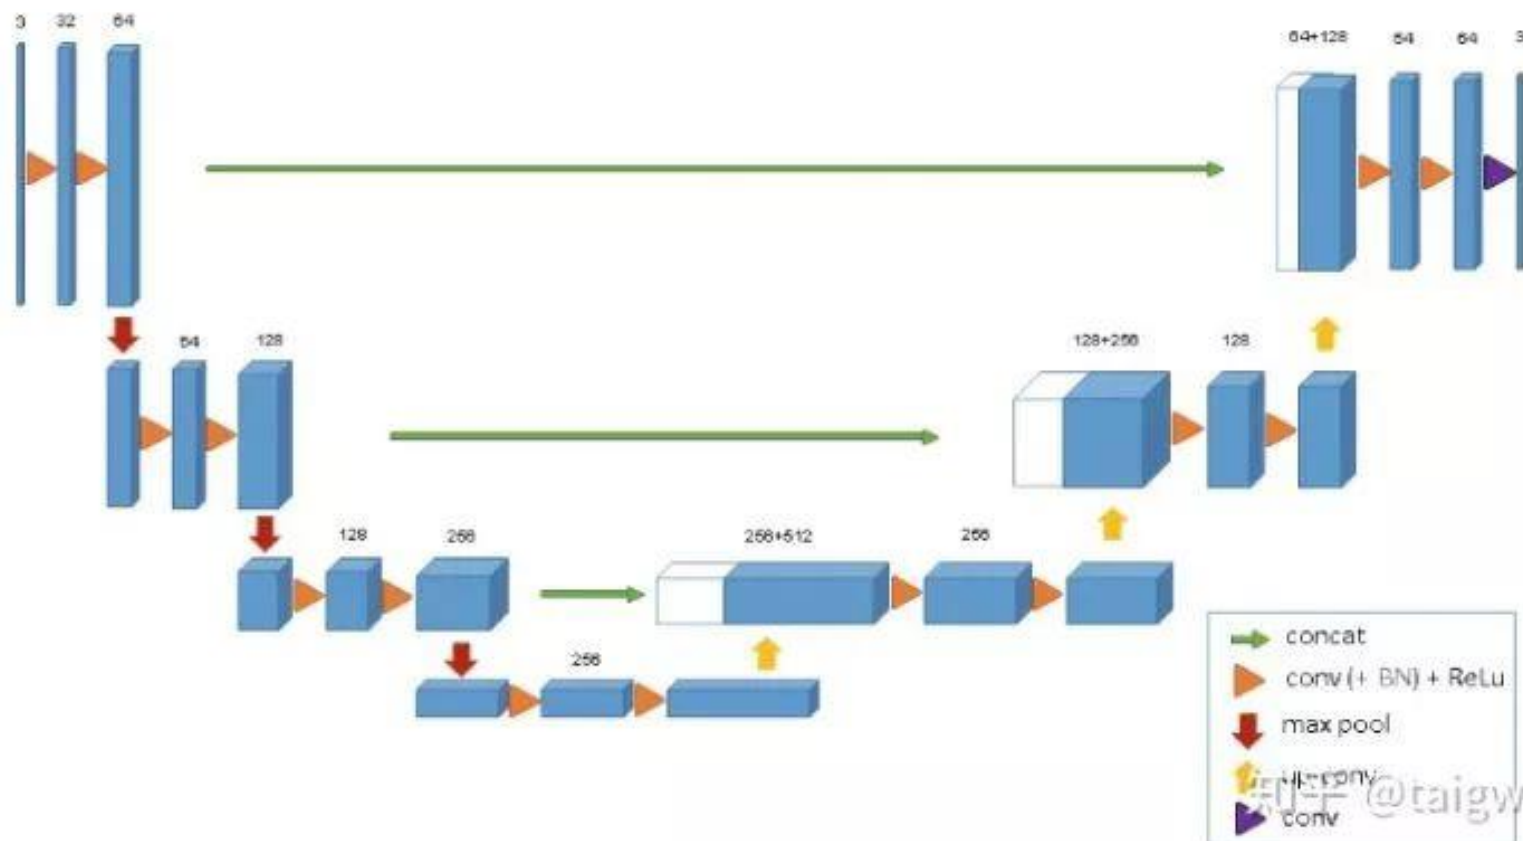


# U-NET

- A special kind of encoder-decoder network
- $N$  encoder layers
- $N$  decoder layers
- Every  $i$ -th encoder layer is connected with  $(N-i)$  decoder layer



# Why U-NET?



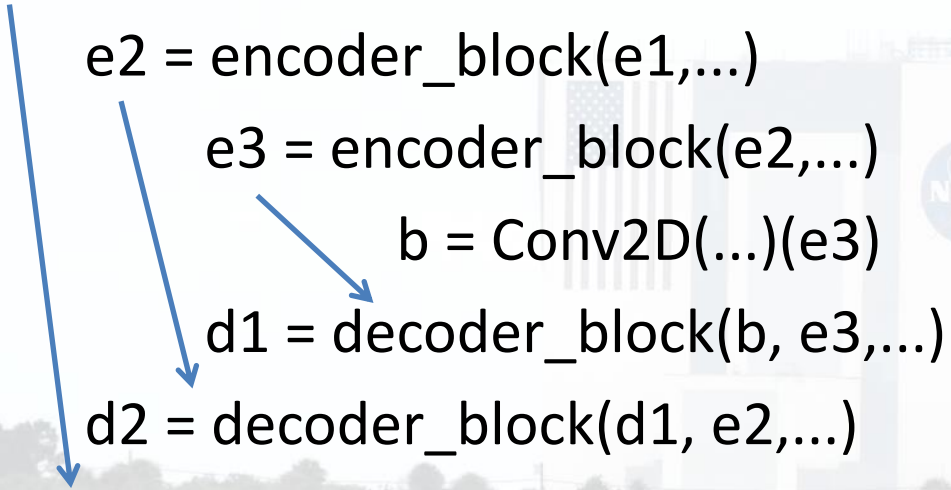
<https://programmer.group/unet-network-magic-changes-those-things.html>

# UNET simplified architecture

```
input = Input(...)
e1 = encoder_block(input, layers, filters,...)
e2 = encoder_block(e1,...)
e3 = encoder_block(e2,...)
b = Conv2D(...)(e3)
d1 = decoder_block(b, e3,...)
d2 = decoder_block(d1, e2,...)
d3 = decoder_block(d2, e1,...)
output = Activation('tanh')(d3)
model = Model(input,output)
```

# UNET simplified architecture

```
input = Input(...)
e1 = encoder_block(input, layers, filters,...)
    e2 = encoder_block(e1,...)
        e3 = encoder_block(e2,...)
            b = Conv2D(...)(e3)
                d1 = decoder_block(b, e3,...)
                    d2 = decoder_block(d1, e2,...)
                        d3 = decoder_block(d2, e1,...)
                            output = Activation('tanh')(d3)
                                model = Model(input,output)
```

A diagram illustrating the UNET architecture. It shows a sequence of encoder blocks (e1, e2, e3) and decoder blocks (d1, d2, d3). Blue arrows indicate skip connections from the output of each encoder block to the corresponding decoder block: from e1 to d3, from e2 to d2, and from e3 to d1. The background features a faint image of a NASA space shuttle launch.



# encoder

```
def encoder_block(layer_in, n_filters, batchnorm=True):  
    g = Conv2D(n_filters, (4,4), strides=(2,2), padding='same',  
              kernel_initializer=init)(layer_in)  
    if batchnorm:  
        g = BatchNormalization()(g, training=True)  
    g = LeakyReLU(alpha=0.2)(g)  
    return g
```

# decoder

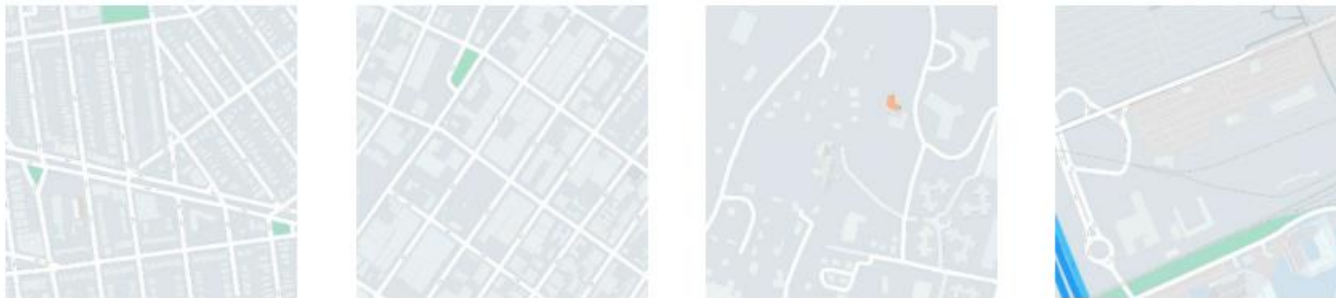
```
def decoder_block(layer_in, skip_in, n_filters, dropout=True):  
    g = Conv2DTranspose(n_filters, (4,4), strides=(2,2),  
        padding='same', kernel_initializer=init)(layer_in)  
    g = BatchNormalization()(g, training=True)  
    if dropout:  
        g = Dropout(0.5)(g, training=True)  
    g = Concatenate()([g, skip_in]) # merge with skip connection  
    g = Activation('relu')(g)  
    return g
```

# Maps with U-NET

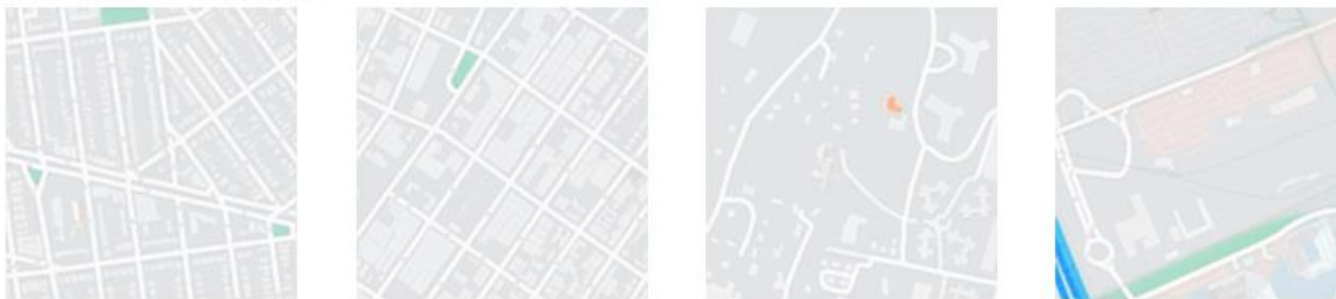
- Good results in few epochs
  - Not surprising – there are direct connections!
- autoencoder\_map\_unet.ipynb***

Epoch 10

Original training images



Reconstructed training images



# Real applications

Some examples:

- Colorization
  - BW image -> color image
- Super-resolution
  - image 64x64 > image 256x256
- Image segmentation:
  - [https://keras.io/examples/vision/oxford\\_pets\\_image\\_segmentation/](https://keras.io/examples/vision/oxford_pets_image_segmentation/)
- Creating analogy
  - satellite image -> map

# Colorization

- Simple CNN network (*colorize.ipynb*)

- $N \times N \times 1 \rightarrow N \times N \times 3$

- Architecture:

```
input_img = Input(shape=image_shape)
```

```
x = Conv2D(filters = 16, kernel_size = (3, 3), activation='relu',  
padding='same')(input_img)
```

```
x = Conv2D(filters = 32, kernel_size = (3, 3), activation='relu',  
padding='same')(x)
```

```
x = Conv2D(filters = 64, kernel_size = (3, 3), activation='relu',  
padding='same')(x)
```

```
output_img = Conv2D(3, (3, 3), padding='same')(x)
```

```
model = Model(input_img, output_img)
```

# Colorization

- Simple CNN network
  - $N \times N \times 1 \rightarrow N \times N \times 3$
- Architecture (simplified notation):

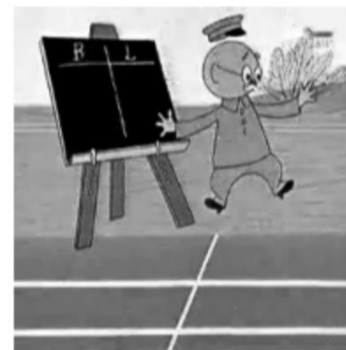
```
args = {"activation": "relu", "padding": "same", "kernel_size": (3,3)}
input_img = Input(shape=image_shape)
x = Conv2D(filters = 16, **args)(input_img)
x = Conv2D(filters = 32, **args)(x)
x = Conv2D(filters = 64, **args)(x)
output_img = Conv2D(3, (3, 3), padding='same')(x)
model = Model(input_img, output_img)
```



# Results

- Not very good...
- ...but not very bad as well!

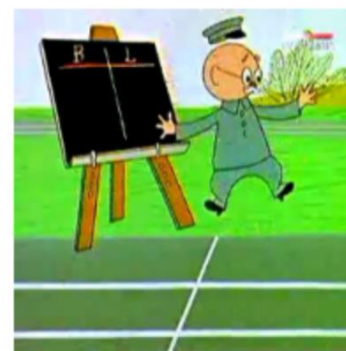
Black&amp;White images



Colorized BW images



Original images



# Important property of CNN

- The number of weights for CNN is independent of the image resolution!
- Conv2D(filters = 16, kernel\_size = (3, 3)) always has
  - $16 * 3 * 3 + 16 = 160$  weights
  - regardless of an image size!
- The next layer Conv2D(32,(3,3)) always has
  - $16 * 32 * 3 * 3 + 32 = 4640$  weights
  - regardless of an image size!
- Pure CNN models work for images with any size!

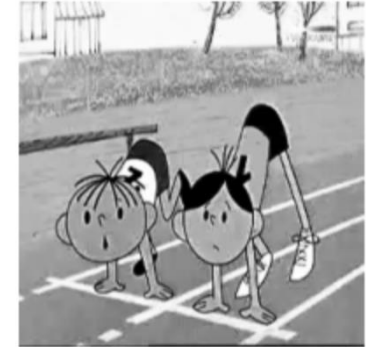
# Using UNET architecture

- Necessary to add two channels to BW images:
  - `bwImages = np.concatenate((bwImages,bwImages,bwImages),axis=3)`
- Results: much better just after few epochs
- ***colorize\_unet.ipynb***

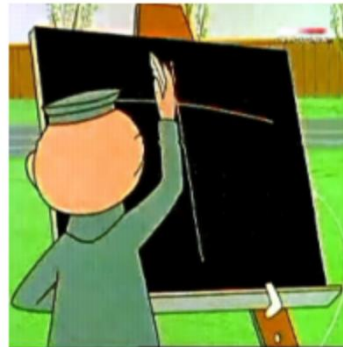
# UNET Results

- Much better!

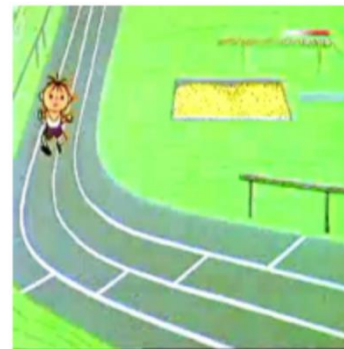
Black&amp;White images



Colorized BW images



Original images





# Super-resolution

- Turn an image with low resolution into the image with high resolution
- The state of the art established during:
  - New Trends in Image Restoration and Enhancement (NTIRE) workshop and challenge on image super-resolution
  - part of the CVPR conference
  - several editions: 2017-2021
- Different possible architectures

# The simplest example

- ***supersampling\_bolek.ipynb***

- 64x64 -> 256x256

- The model:

```
conv_args = {"activation": "relu", "padding": "same", }
```

```
inputs = Input(shape=image_shape)
```

```
x = Conv2D(64, 5, **conv_args)(inputs)
```

```
x = Conv2D(64, 3, **conv_args)(x)
```

```
x = Conv2D(32, 3, **conv_args)(x)
```

```
x = Conv2D(channels * (up_factor ** 2), 3, **conv_args)(x)
```

```
outputs = tf.nn.depth_to_space(x, up_factor)
```

```
model = Model(inputs, outputs)
```



# depth\_to\_space

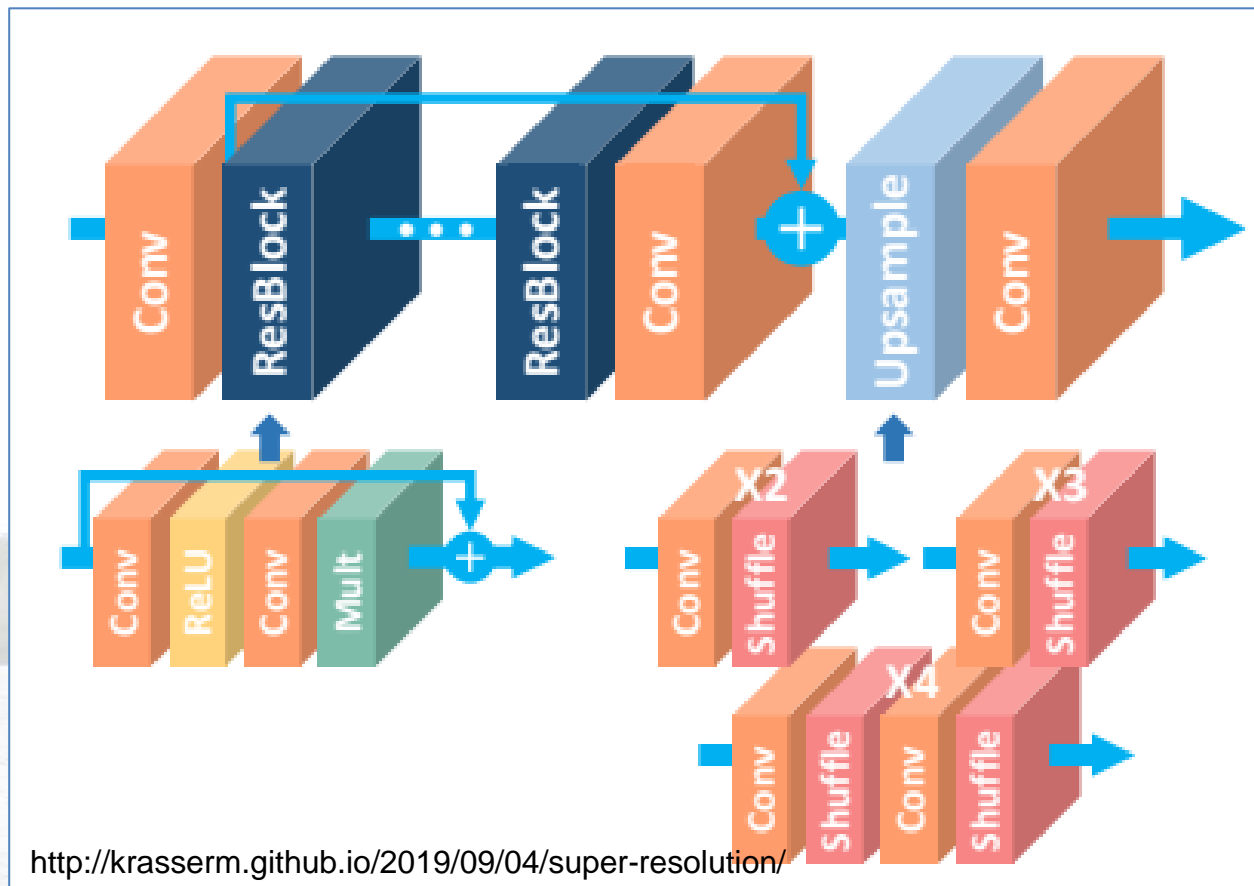
- Conversion with scale factor:  $s$
- General task:
  - $(W, H, C) > (s*W, s*H, C)$
- Depth to space layer:
  - $(W, H, C*s^2) > (s*W, s*H, C)$
- Example:
  - $(32, 32, 3)$
  - ...
  - $(32, 32, 3*4^2)$
  - depth\_to\_space layer
  - $(32*4, 32*4, 3)$

# More sophisticated architectures

- Enhanced Deep Residual Networks for Single Image Super-Resolution (EDSR)
  - winner of NTIRE 2017
- Wide Activation for Efficient and Accurate Image Super-Resolution (WDSR)
  - winner of NTIRE 2018
- Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network (SRGAN)
  - GAN network

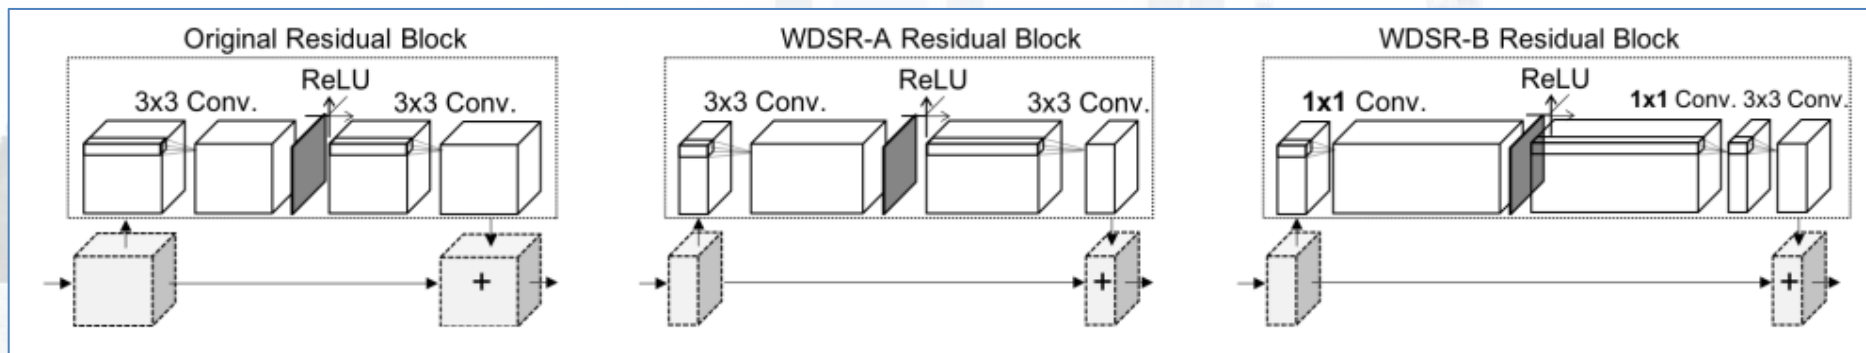
# EDSR

- Residual network with Conv2D-RELU-Conv2D-Mult blocks



# WDSR

- Extension of EDSR:
  - increases the number of channels in residual blocks
  - reduces the number of channels in mapping path
  - the same number of weights



<http://krasserm.github.io/2019/09/04/super-resolution/>

# Evaluation

- How to evaluate the correctness of superscaling?
- The obvious idea: calculate the difference between the generated image and the real image:
  - L2 norm
  - L1 norm
  - Binary crossentropy
- Problem: images perceived as blurred have typically good results
- A step forward:
  - use the additional network (discriminator!) to judge the correctness!
- SRGAN

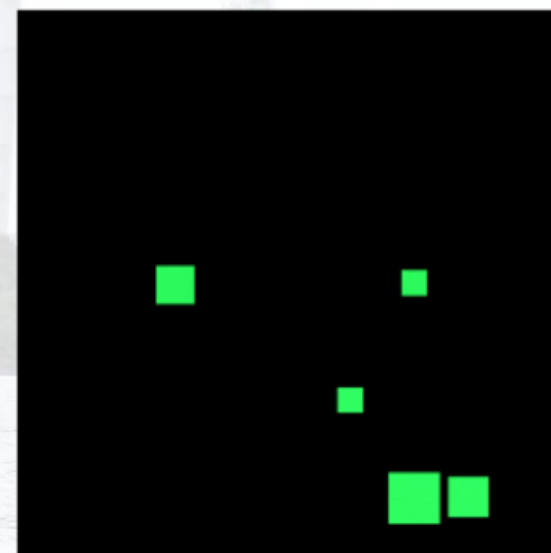
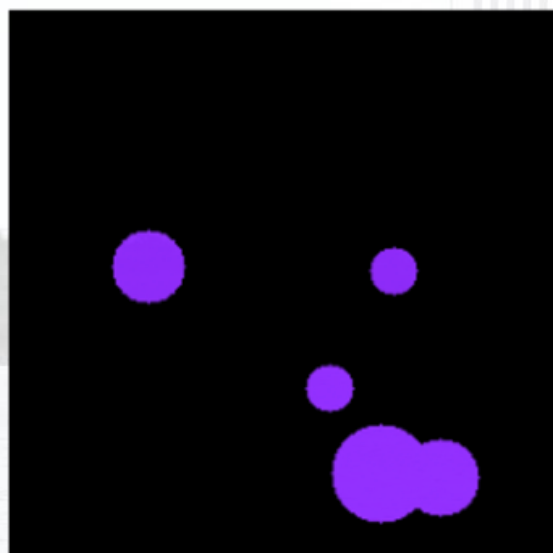
# Ready to use

- More information:
  - <http://krasserm.github.io/2019/09/04/super-resolution/>
- Library with code:
  - <https://github.com/krasserm/super-resolution>
- Execution examples
  - article.ipynb
  - example-esdr.ipynb
  - example-wdsr.ipynb
  - example-srgan.ipynb



# Image to image (pix2pix)

- GAN that converts one image to another
- Input and output images are different but there is **analogy** between them
- A simple example: turn violet circles to green rectangles

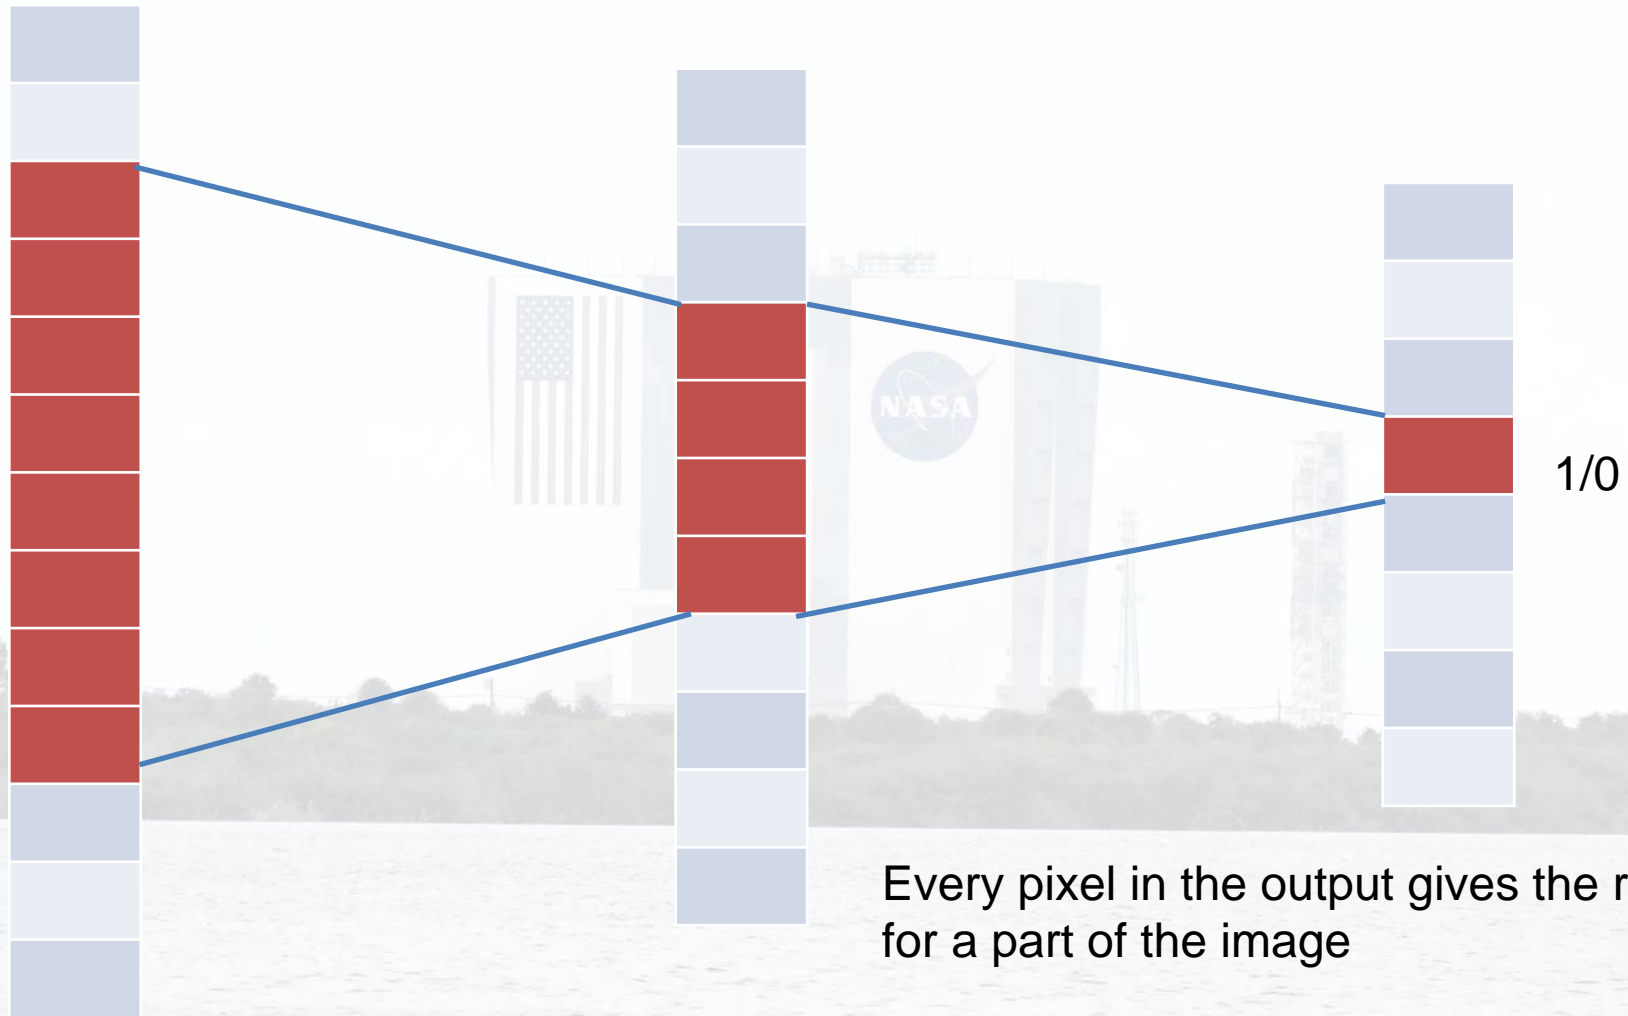


# Working example

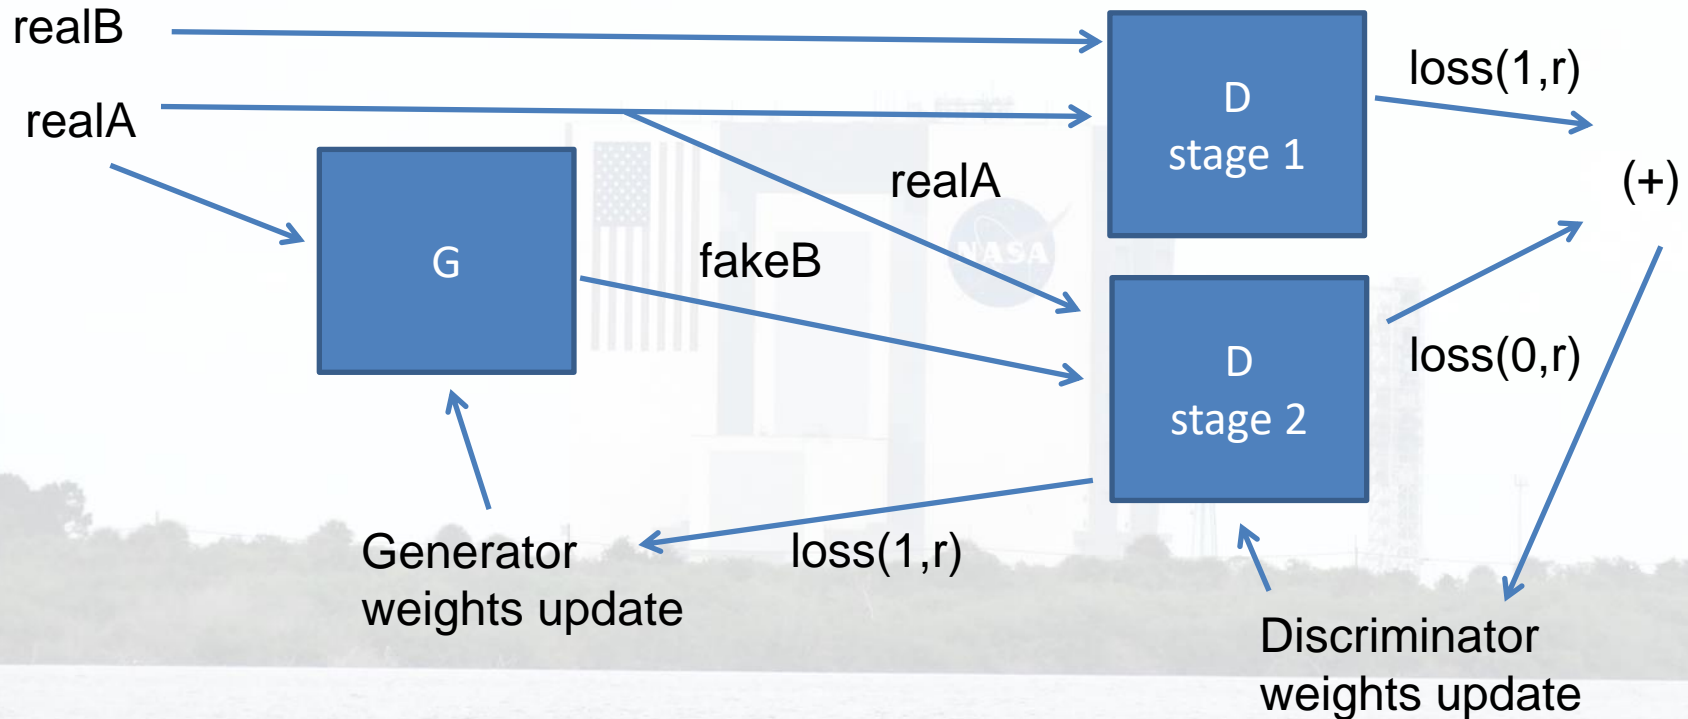
## ***pix2pix.ipynb***

- generator:
  - UNET network (encoder-decoder with residuals)
- discriminator:
  - PatchGAN
  - It does not return one value 1/0
  - It returns a matrix of 1/0 values
  - Every pixel in the matrix refers to some part of the image
    - the parts overlap!
  - The architecture works with images of any size!

# PatchGAN



# Pix2Pix GAN architecture



# A bit different GAN creation

- This time we don't use the GradientTape!
  - discriminator\_model – will be trained by itself
  - generator\_model – will be trained through the "gan\_model"
- Creation of gan\_model (the network used to train generator)

```
def gan_model(generator_model, discriminator_model):
```

```
    for layer in discriminator_model.layers: # discriminator will not be trained
```

```
        layer.trainable = False
```

```
    input_src = Input(shape=image_shape)
```

```
    gen_out = generator_model(input_src)
```

```
    disc_out = discriminator_model([input_src, gen_out])
```

```
    model = Model(input_src, [disc_out, gen_out])
```

```
    return model
```

# Training GAN

- Preparation:

```
patch = d_model.output_shape[1] # output of discriminator
```

```
steps = int(len(trainImgs) / batch) # steps per epoch
```

```
all_ones = np.ones((batch, patch, patch, 1)) # expected output for real
```

```
all_zeros = np.zeros((batch, patch, patch, 1)) # expected output for fake
```

- One learning step:

```
for epoch in range(epochs):
```

```
    for i in range(steps):
```

```
        realA, realB = generate_real_samples(batch)
```

```
        fakeB = g_model.predict(realA)
```

```
        d_loss1 = d_model.train_on_batch([realA, realB], all_ones )
```

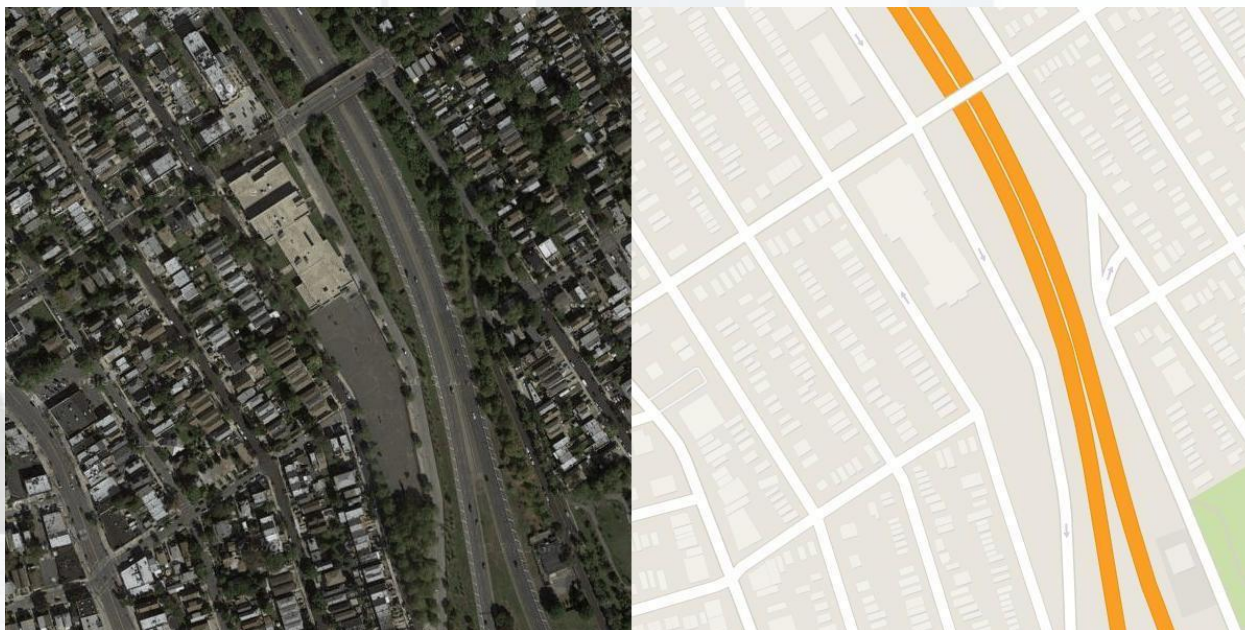
```
        d_loss2 = d_model.train_on_batch([realA, fakeB], all_zeros)
```

```
        g_loss, _, _ = gan_model.train_on_batch(realA, [all_ones, realB])
```



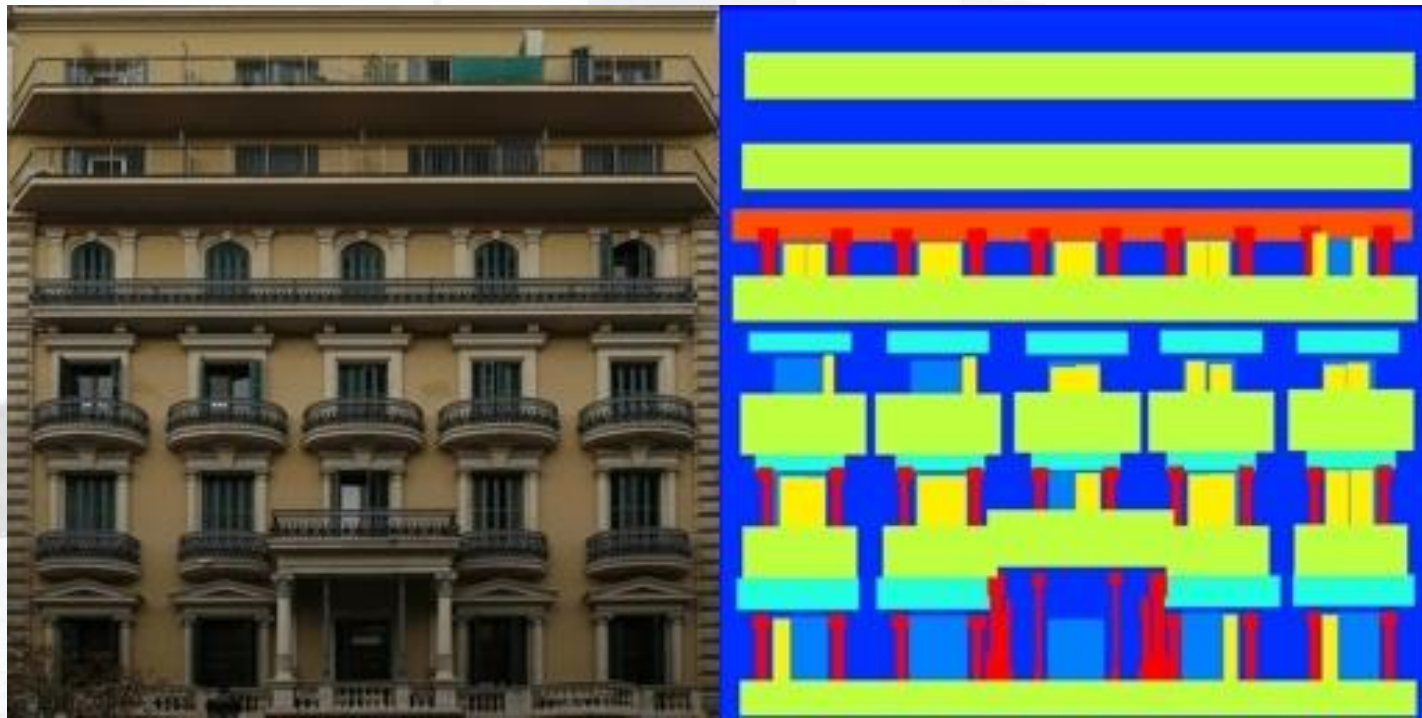
# map2image example

- Load pairs: satellite image and google map
- ***pix2pix\_map.ipynb***



# Ready-to-use solution

- <https://github.com/affinelayer/pix2pix-tensorflow>
- Dataset preparation: set of images side by side:



# Using pix2pix

- Training on facades:
  - `python pix2pix.py --mode train --output_dir facades_train --max_epochs 200 --input_dir facades/train --which_direction BtoA`
- Testing facades:
  - `python pix2pix.py --mode test --output_dir facades_test --input_dir facades/val --checkpoint facades_train`
- Result: the html file with pairs of images

# Available datasets

- <https://www.github.com/affinelayer/pix2pix-tensorflow-models.git> static/models
  - facades
  - edges2cats
  - edges2shoes
  - edges2handbags
- Online example:
  - <https://affinelayer.com/pixsrv/>
- It is possible to start your own server:
  - cd server
  - `serve.py --port 8001`

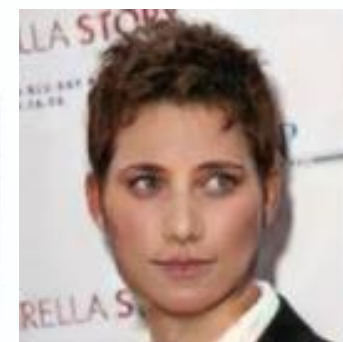


# Using ready-to-use models

- Install tensorflow\_examples package
  - `pip install git+https://github.com/tensorflow/examples.git`
- Import the package:  
`from tensorflow_examples.models.pix2pix import pix2pix`
- Create the generator and discriminator:  
`generator = pix2pix.unet_generator(...)`  
`discriminator = pix2pix.discriminator(...)`

# Problems with pix2pix

- It requires pairs of analogous images
- It is not always possible
- Simple example: change man face to woman face



- Problem: we don't have many pairs like that
- If we had a software producing such pairs we would not need any GAN!



# CycleGAN

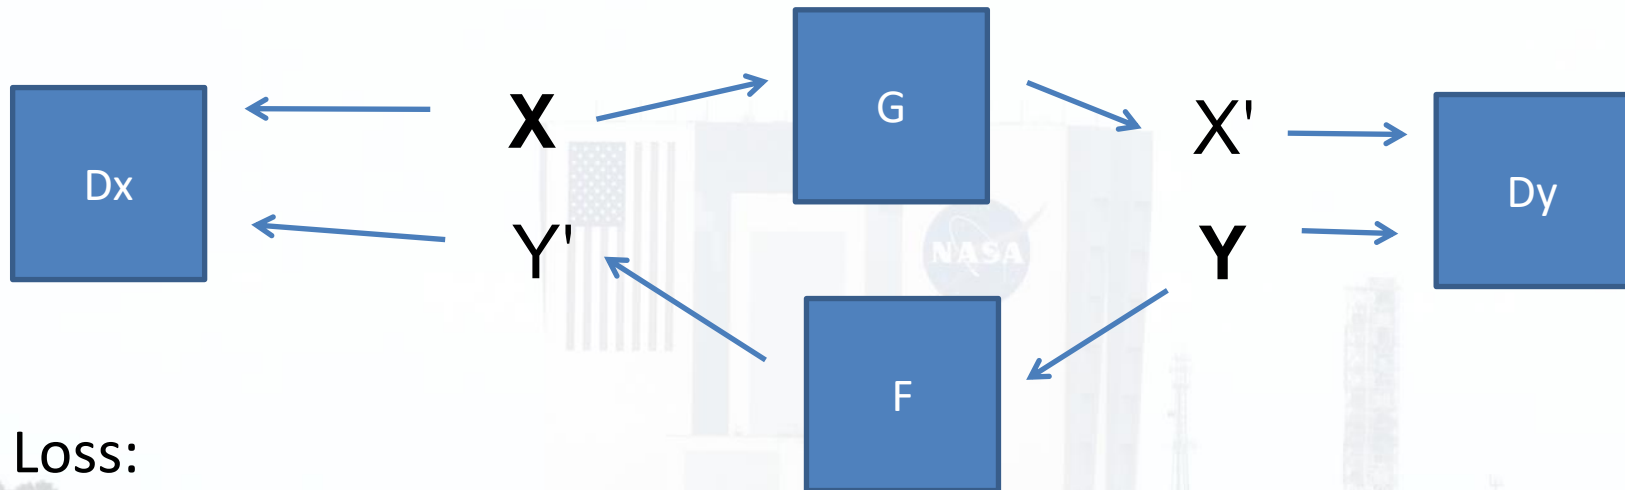
- Instead of preparing pairs of images we prepare sets of images
  - without one-to-one relationships!
- For example:
  - set of female images (X)
  - set of male images (Y)
- We train the network to generate images based on X that look like Y
- ...and images based on Y that look like X (a cycle!)

# CycleGAN architecture

- Two generators:  $G$  and  $F$ 
  - $G$  translates  $X$  to  $Y$
  - $F$  translates  $Y$  to  $X$
- Two discriminators  $D_x$  and  $D_y$ :
  - $D_x$  checks if  $X$  is genuine or fake
  - $D_y$  checks if  $Y$  is genuine or fake

# CycleGAN architecture

Two sets of images: X and Y



Loss:

$$Dx\_loss = bce(X,1) + bce(Y',0)$$

$$Dy\_loss = bce(Y,1) + bce(X',0)$$

$$G\_loss = bce(X',1) + \dots$$

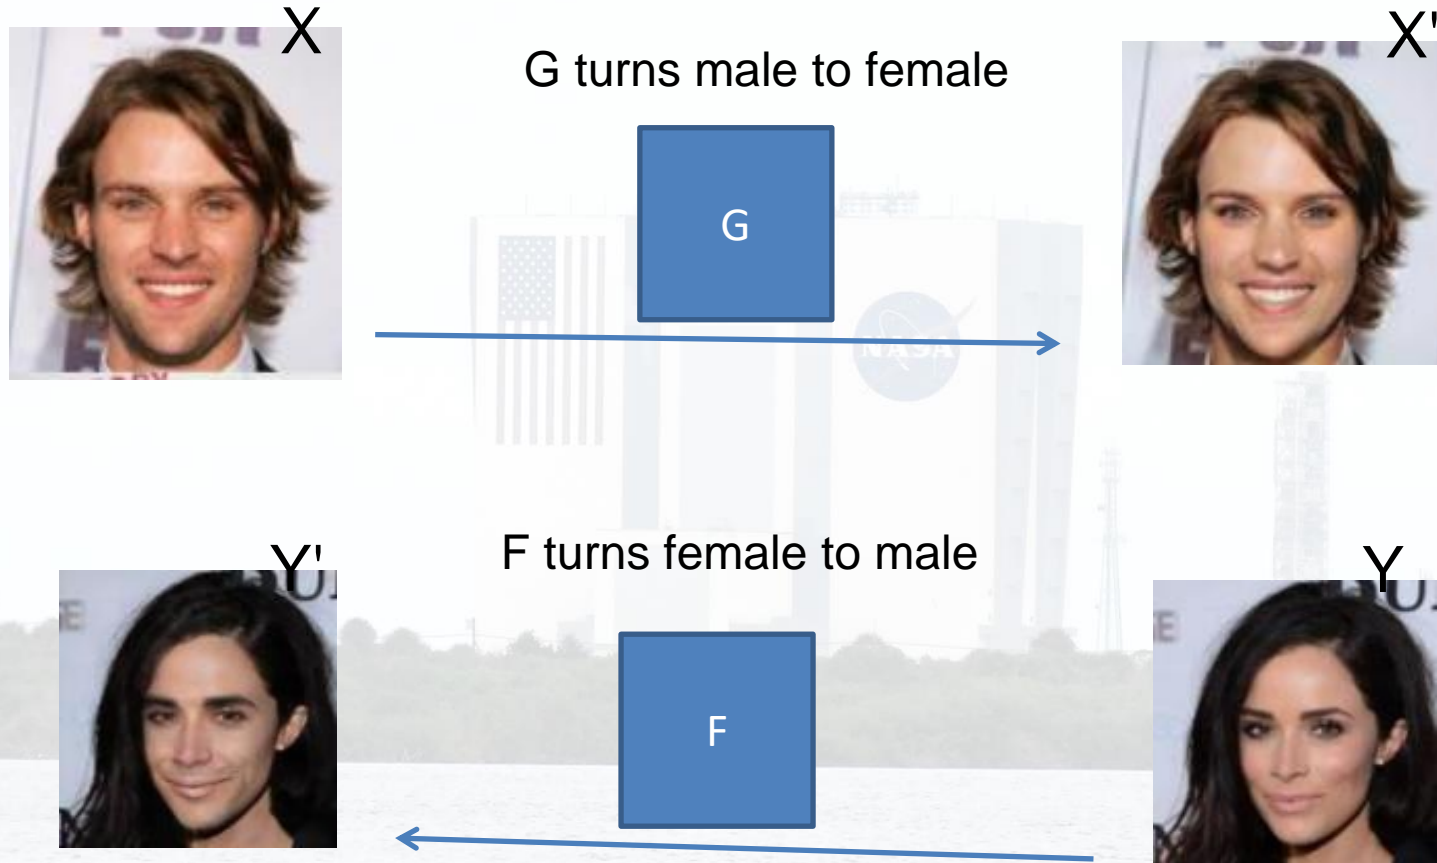
$$F\_loss = bce(Y',1) + \dots$$

bce = binary crossentropy

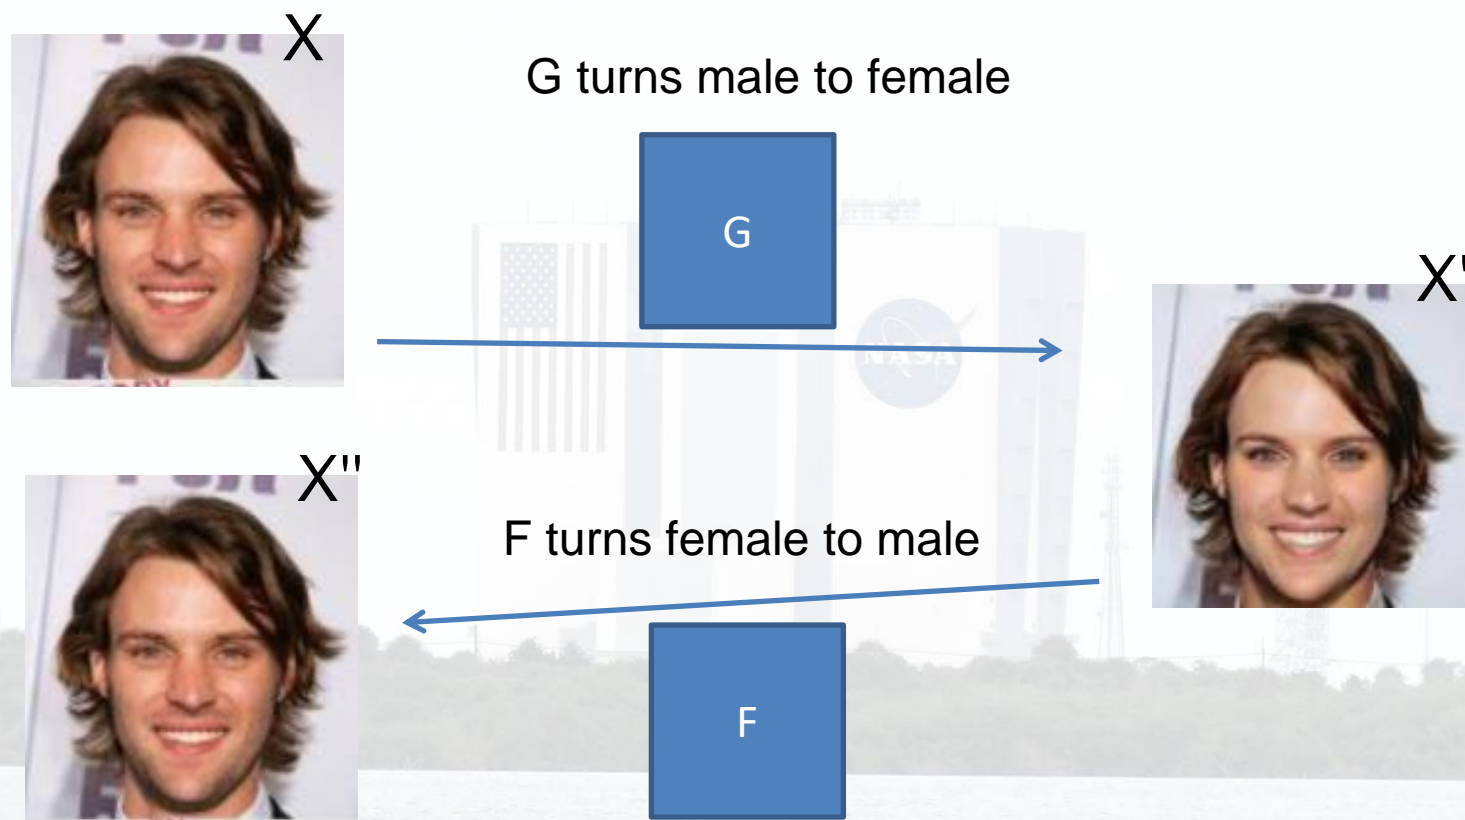
# Two additional losses for CycleGAN

- **Cycle loss:** after the cycle the image should look the same
  - $X' = G(X)$
  - $X'' = F(X')$
  - $\text{cycle\_loss\_x} = |X'' - X| = |F(G(X)) - X|$
  - $\text{cycle\_loss\_y} = |G(F(Y)) - Y|$
  - $\text{total\_cycle\_loss} = \text{cycle\_loss\_x} + \text{cycle\_loss\_y}$
- **Identity loss:** after the "reverse generation" the image should look the same
  - $\text{identity\_loss\_x} = |F(X) - X|$
  - $\text{identity\_loss\_y} = |G(Y) - Y|$

# For our example



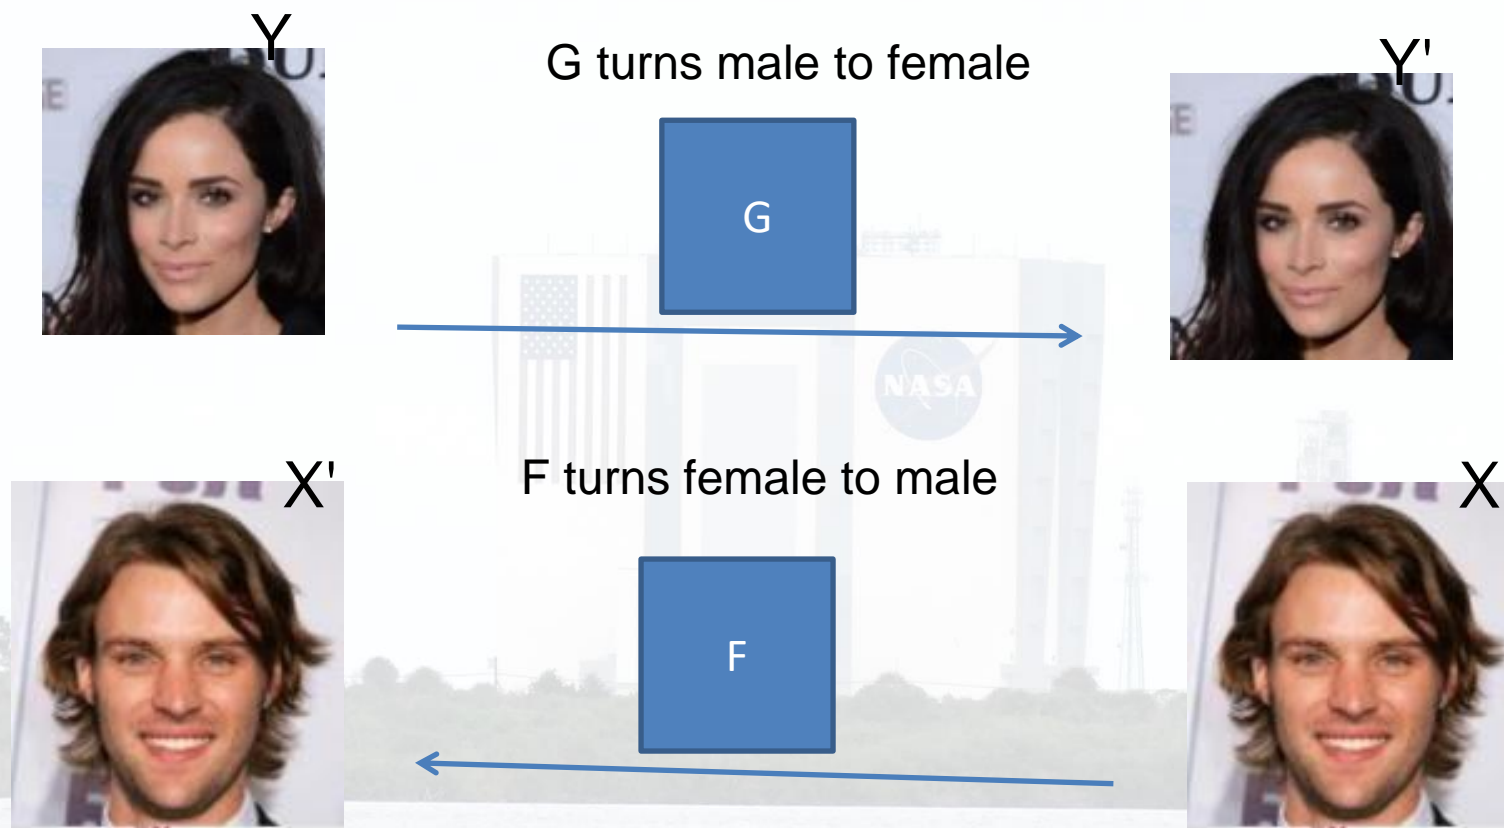
# Cycle loss



$X''$  should be similar to  $X$ :  $\text{cycle\_loss\_x} = |X - X''| = |X - F(G(X))|$



# Identity loss



Turning male image to male should have no effect:  $\text{identity\_loss\_x} = |X - F(X)|$

# Rules for loss calculation

- Discriminator X should recognize male faces
- Discriminator Y should recognize female faces
- Cycle:
  - Male face after generator G should turn to female, and this changed image after generator F should turn to male again – the same as at the beginning
  - Female face after generator F should turn to male, and this changed image after generator G should turn to female again
- Identity:
  - Male face used as input to generator F should remain male
  - Female face used as input to generator G should remain female

# One step

input: real\_x, real\_y

## **# generate images**

fake\_y = generator\_g(real\_x, training=True)

cycled\_x = generator\_f(fake\_y, training=True)

fake\_x = generator\_f(real\_y, training=True)

cycled\_y = generator\_g(fake\_x, training=True)

same\_x = generator\_f(real\_x, training=True)

same\_y = generator\_g(real\_y, training=True)

## **# check results**

disc\_real\_x = discriminator\_x(real\_x, training=True)

disc\_real\_y = discriminator\_y(real\_y, training=True)

disc\_fake\_x = discriminator\_x(fake\_x, training=True)

disc\_fake\_y = discriminator\_y(fake\_y, training=True)

# Calculate loss

*bce* – binary cross entropy, *abs* – mean absolute error

## # discriminators losses

```
disc_x_loss = bce([1],disc_real_x) + bce([0],disc_fake_x)
```

```
disc_y_loss = bce([1],disc_real_y) + bce([0], disc_fake_y)
```

## # generators losses

```
gen_g_loss = bce([1],disc_fake_y)
```

```
gen_f_loss = bce([1],disc_fake_x)
```

```
total_cycle_loss = abs(real_x, cycled_x) + abs(real_y, cycled_y)
```

```
identity_loss_x = abs(real_x, same_x)
```

```
identity_loss_y = abs(real_y, same_y)
```

## # total generator losses

```
total_gen_g_loss = gen_g_loss + total_cycle_loss + identity_loss_y
```

```
total_gen_f_loss = gen_f_loss + total_cycle_loss + identity_loss_x
```

# Apply gradients

```
# Calculate the gradients for generators and discriminators
```

```
g_grads = tape.gradient(total_gen_g_loss, generator_g.trainable_variables)
```

```
f_grads = tape.gradient(total_gen_f_loss, generator_f.trainable_variables)
```

```
dx_grads = tape.gradient(disc_x_loss, discriminator_x.trainable_variables)
```

```
dy_grads = tape.gradient(disc_y_loss, discriminator_y.trainable_variables)
```

```
# Apply the gradients to the networks
```

```
g_optimizer.apply_gradients(zip(g_grads, generator_g.trainable_variables))
```

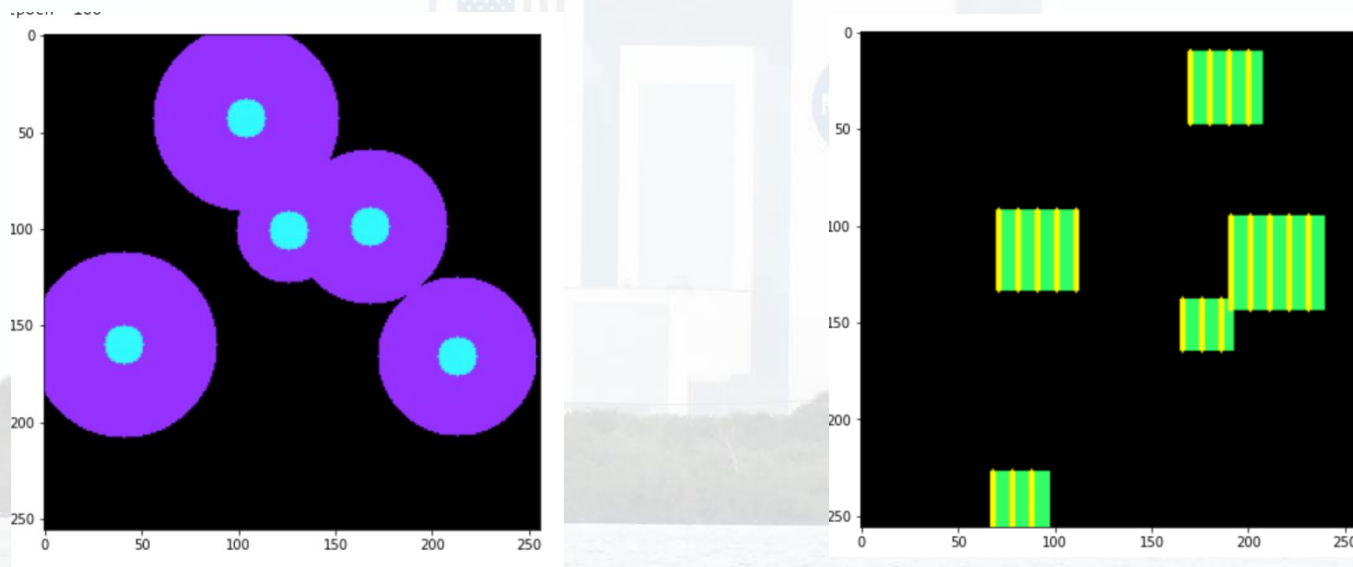
```
f_optimizer.apply_gradients(zip(f_grads, generator_f.trainable_variables))
```

```
dx_opt.apply_gradients(zip(dx_grads, discriminator_x.trainable_variables))
```

```
dy_opt.apply_gradients(zip(dy_grads, discriminator_y.trainable_variables))
```

# CycleGAN example

- *cyclegan.ipynb*
- Changing circles to squares





# A classic example

- Notebook from the Tensorflow tutorial
  - <https://www.tensorflow.org/tutorials/generative/cyclegan>
- Changing horses to zebras

Input Image



Predicted Image



# Other datasets

- [https://people.eecs.berkeley.edu/~taesung\\_park/CycleGAN/datasets/](https://people.eecs.berkeley.edu/~taesung_park/CycleGAN/datasets/)
- apple2orange.zip
- cezanne2photo.zip
- iphone2dslr\_flower.zip
- monet2photo.zip
- summer2winter\_yosemite.zip
- vangogh2photo.zip
- ...

# Summary

- Autoencoders and U-Networks
  - may be used for image conversion (denoising, colorization, supersampling,...)
- Pix2pix
  - converts one image to another
  - we need pairs of images
- CycleGAN
  - builds generators that convert one type of images into another type
- There are a lot of interesting applications!

# Deep Learning in Python

Next lecture: Object detection



## Autoencoders and Pix2pix networks

Paweł Kasprowski, PhD, DSc.